

DAI Assignment-1 Report

Task: Adversarial Perturbations: Using CIFAR10 dataset and protocol

Objectives:

- **implement FGSM on your own and using a basic CNN architecture with three convolution layers, you need to show the impact of the attack.**
- **Use two deep learning architectures of your choice, PGD attack and one other than FGSM.**
- **Using the SVHN dataset, you need to show the impact of the attacks and compare and contrast them**

Procedure:

- Import required packages. (PyTorch, NumPy, Matplotlib, torchmetrics ...etc.)
- Setting up device-agnostic code for faster training.
- Downloading the Dataset using PyTorch vision datasets.
- Visualizing random data samples from the dataset.
- Making compose transform to convert image to tensors
- Converting the data to torch. Tensors



- Making data loaders from the datasets.
- Defining Basic CNN-Architecture Class Definition.
- Defining training and testing steps for training method.
- Creating the model object defined above, and printing model summary.

=====				
Layer (type (var_name))	Input Shape	Output Shape	Param #	
Trainable				
=====				
BasicCNN (BasicCNN)	[32, 3, 32, 32]	[32, 10]	--	True
└ Sequential (conv_block1)	[32, 3, 32, 32]	[32, 10, 16, 16]	--	True
└ Conv2d (0)	[32, 3, 32, 32]	[32, 10, 32, 32]	280	True
└ ReLU (1)	[32, 10, 32, 32]	[32, 10, 32, 32]	--	--
└ Conv2d (2)	[32, 10, 32, 32]	[32, 10, 32, 32]	910	True
└ ReLU (3)	[32, 10, 32, 32]	[32, 10, 32, 32]	--	--
└ MaxPool2d (4)	[32, 10, 32, 32]	[32, 10, 16, 16]	--	--
└ Sequential (conv_block2)	[32, 10, 16, 16]	[32, 10, 8, 8]	--	True
└ Conv2d (0)	[32, 10, 16, 16]	[32, 10, 16, 16]	910	True
└ ReLU (1)	[32, 10, 16, 16]	[32, 10, 16, 16]	--	--
└ MaxPool2d (2)	[32, 10, 16, 16]	[32, 10, 8, 8]	--	--
└ Sequential (classifier)	[32, 10, 8, 8]	[32, 10]	--	True
└ Flatten (0)	[32, 10, 8, 8]	[32, 640]	--	--
└ Linear (1)	[32, 640]	[32, 10]	6,410	True
└ ReLU (2)	[32, 10]	[32, 10]	--	--
└ Linear (3)	[32, 10]	[32, 10]	110	True
=====				
Total params: 8,620				
Trainable params: 8,620				
Non-trainable params: 0				
Total mult-adds (M): 46.66				
=====				
Input size (MB): 0.39				
Forward/backward pass size (MB): 5.90				
Params size (MB): 0.03				
Estimated Total Size (MB): 6.33				
=====				


- Training the model with cross CrossEntropyLoss and Adam optimizer.

- Training Results:

Epoch: 1 | train_loss: 1.8998 | train_acc: 0.2903 | test_loss: 1.6249 | test_acc: 0.3900
Epoch: 2 | train_loss: 1.5543 | train_acc: 0.4167 | test_loss: 1.4560 | test_acc: 0.4466
Epoch: 3 | train_loss: 1.4405 | train_acc: 0.4552 | test_loss: 1.4071 | test_acc: 0.4616
Epoch: 4 | train_loss: 1.3858 | train_acc: 0.4759 | test_loss: 1.3644 | test_acc: 0.4818
Epoch: 5 | train_loss: 1.3429 | train_acc: 0.4948 | test_loss: 1.3139 | test_acc: 0.4952
Epoch: 6 | train_loss: 1.2986 | train_acc: 0.5111 | test_loss: 1.2874 | test_acc: 0.5075
Epoch: 7 | train_loss: 1.2697 | train_acc: 0.5188 | test_loss: 1.2707 | test_acc: 0.5210
Epoch: 8 | train_loss: 1.2446 | train_acc: 0.5301 | test_loss: 1.2369 | test_acc: 0.5285
Epoch: 9 | train_loss: 1.2282 | train_acc: 0.5334 | test_loss: 1.2266 | test_acc: 0.5294
Epoch: 10 | train_loss: 1.2102 | train_acc: 0.5395 | test_loss: 1.2583 | test_acc: 0.5237
Epoch: 11 | train_loss: 1.1976 | train_acc: 0.5455 | test_loss: 1.2013 | test_acc: 0.5384
Epoch: 12 | train_loss: 1.1851 | train_acc: 0.5503 | test_loss: 1.1918 | test_acc: 0.5469
Epoch: 13 | train_loss: 1.1714 | train_acc: 0.5554 | test_loss: 1.1942 | test_acc: 0.5437
Epoch: 14 | train_loss: 1.1607 | train_acc: 0.5604 | test_loss: 1.1757 | test_acc: 0.5521
Epoch: 15 | train_loss: 1.1471 | train_acc: 0.5625 | test_loss: 1.1699 | test_acc: 0.5525
Epoch: 16 | train_loss: 1.1384 | train_acc: 0.5688 | test_loss: 1.1537 | test_acc: 0.5535
Epoch: 17 | train_loss: 1.1294 | train_acc: 0.5722 | test_loss: 1.1575 | test_acc: 0.5516
Epoch: 18 | train_loss: 1.1216 | train_acc: 0.5728 | test_loss: 1.1542 | test_acc: 0.5586
Epoch: 19 | train_loss: 1.1125 | train_acc: 0.5758 | test_loss: 1.1549 | test_acc: 0.5571
Epoch: 20 | train_loss: 1.1035 | train_acc: 0.5834 | test_loss: 1.1800 | test_acc: 0.5536
Epoch: 21 | train_loss: 1.0944 | train_acc: 0.5838 | test_loss: 1.1302 | test_acc: 0.5666
Epoch: 22 | train_loss: 1.0907 | train_acc: 0.5882 | test_loss: 1.1331 | test_acc: 0.5654
Epoch: 23 | train_loss: 1.0824 | train_acc: 0.5877 | test_loss: 1.1142 | test_acc: 0.5723
Epoch: 24 | train_loss: 1.0765 | train_acc: 0.5933 | test_loss: 1.1337 | test_acc: 0.5670
Epoch: 25 | train_loss: 1.0707 | train_acc: 0.5953 | test_loss: 1.1300 | test_acc: 0.5707
Epoch: 26 | train_loss: 1.0674 | train_acc: 0.5951 | test_loss: 1.1253 | test_acc: 0.5702
Epoch: 27 | train_loss: 1.0604 | train_acc: 0.5968 | test_loss: 1.1083 | test_acc: 0.5770
Epoch: 28 | train_loss: 1.0573 | train_acc: 0.6010 | test_loss: 1.1062 | test_acc: 0.5779
Epoch: 29 | train_loss: 1.0530 | train_acc: 0.5994 | test_loss: 1.1077 | test_acc: 0.5777
Epoch: 30 | train_loss: 1.0495 | train_acc: 0.5976 | test_loss: 1.1067 | test_acc: 0.5780
Epoch: 31 | train_loss: 1.0444 | train_acc: 0.6038 | test_loss: 1.1144 | test_acc: 0.5767
Epoch: 32 | train_loss: 1.0432 | train_acc: 0.6019 | test_loss: 1.1112 | test_acc: 0.5772
Epoch: 33 | train_loss: 1.0402 | train_acc: 0.6042 | test_loss: 1.1056 | test_acc: 0.5822
Epoch: 34 | train_loss: 1.0329 | train_acc: 0.6067 | test_loss: 1.1055 | test_acc: 0.5805
Epoch: 35 | train_loss: 1.0345 | train_acc: 0.6074 | test_loss: 1.0981 | test_acc: 0.5806
Epoch: 36 | train_loss: 1.0319 | train_acc: 0.6094 | test_loss: 1.0967 | test_acc: 0.5817
Epoch: 37 | train_loss: 1.0298 | train_acc: 0.6102 | test_loss: 1.0914 | test_acc: 0.5850
Epoch: 38 | train_loss: 1.0262 | train_acc: 0.6078 | test_loss: 1.1356 | test_acc: 0.5778
Epoch: 39 | train_loss: 1.0236 | train_acc: 0.6122 | test_loss: 1.0860 | test_acc: 0.5898
Epoch: 40 | train_loss: 1.0203 | train_acc: 0.6134 | test_loss: 1.1033 | test_acc: 0.5823
Epoch: 41 | train_loss: 1.0221 | train_acc: 0.6105 | test_loss: 1.1150 | test_acc: 0.5795
Epoch: 42 | train_loss: 1.0172 | train_acc: 0.6086 | test_loss: 1.0864 | test_acc: 0.5877
Epoch: 43 | train_loss: 1.0137 | train_acc: 0.6146 | test_loss: 1.0935 | test_acc: 0.5849

Epoch: 44 | train_loss: 1.0148 | train_acc: 0.6135 | test_loss: 1.1049 | test_acc: 0.5827
Epoch: 45 | train_loss: 1.0129 | train_acc: 0.6167 | test_loss: 1.1177 | test_acc: 0.5830
Epoch: 46 | train_loss: 1.0106 | train_acc: 0.6140 | test_loss: 1.0872 | test_acc: 0.5895
Epoch: 47 | train_loss: 1.0101 | train_acc: 0.6161 | test_loss: 1.0941 | test_acc: 0.5884
Epoch: 48 | train_loss: 1.0068 | train_acc: 0.6158 | test_loss: 1.0885 | test_acc: 0.5867
Epoch: 49 | train_loss: 1.0068 | train_acc: 0.6139 | test_loss: 1.0830 | test_acc: 0.5875
Epoch: 50 | train_loss: 1.0053 | train_acc: 0.6195 | test_loss: 1.0692 | test_acc: 0.6029
total training time: 603.354 sec.

- Saving the trained model, loading, and re-evaluating the training.
- test_loss: 1.0692 | test_acc: 0.6029



- Now, once the model has been trained begin, Attacking the Model.
- Setting epsilon for FGSM

$$perturbed_image = image + epsilon * sign(data_grad) = x + \epsilon * sign(\nabla_x J(\theta, x, y))$$

Finally, in order to maintain the original range of the data, the perturbed image is clipped to range $[0,1]$.

- Defining the attack function.

Python

FGSM attack

```
def fgsm_attack(image: torch.Tensor, epsilon: torch.Tensor,
data_grad: torch.Tensor) -> torch.Tensor:
    # Collect the element-wise sign of the data gradient
    sign_data_grad = data_grad.sign()
    # Create the perturbed image by adjusting each pixel of the
input image
    perturbed_image = image + epsilon*sign_data_grad
    # Adding clipping to maintain [0,1] range
    perturbed_image = torch.clamp(perturbed_image, 0, 1)
    # Return the perturbed image
    return perturbed_image
```

- Testing the attack on the trained model, with different epsilons values.
- Testing Results:

Attacked Examples: 3752

test_loss: 1.0692 | test_acc: 0.6029 | epsilon: 0.0000

Attacked Examples: 9682

test_loss: 1.0692 | test_acc: 0.0310 | epsilon: 0.0625

Attacked Examples: 9466

test_loss: 1.0692 | test_acc: 0.0521 | epsilon: 0.1250

Attacked Examples: 9336

test_loss: 1.0692 | test_acc: 0.0641 | epsilon: 0.1875

Attacked Examples: 9296

test_loss: 1.0692 | test_acc: 0.0680 | epsilon: 0.2500

Attacked Examples: 9296

test_loss: 1.0692 | test_acc: 0.0689 | epsilon: 0.3125


Attacked Examples: 9265

test_loss: 1.0692 | test_acc: 0.0730 | epsilon: 0.3750

Attacked Examples: 9249

test_loss: 1.0692 | test_acc: 0.0737 | epsilon: 0.4375

- Ploting attacked Examples.



Procedure For task 2 (Attack 2 Arch.) :

- Import required packages. (PyTorch, NumPy, Matplotlib, torchmetrics ...etc.)
- Setting up device-agnostic code for faster training.
- Downloading the Dataset using PyTorch vision datasets.
- Visualizing random data samples from the dataset.
- Making compose transform to convert image to tensors
- Converting the data to torch. Tensors



- Making data loaders from the datasets.
- Using torchvision models EfficientNet_B0 & ResNet34 for the task with pre-trained weights and pre-defined transforms.
- Model summaries for both the models are:

```
=====
```

Layer (type (var_name))	Input Shape	Output Shape	Param #
Trainable			
=====			
EfficientNet (EfficientNet)	[32, 3, 224, 224]	[32, 10]	--
Partial			
└─Sequential (features)	[32, 3, 224, 224]	[32, 1280, 7, 7]	--
False			
└─┬─Conv2dNormActivation (0)	[32, 3, 224, 224]	[32, 32, 112, 112]	--
False			
└─└─┬─Conv2d (0)	[32, 3, 224, 224]	[32, 32, 112, 112]	(864)
False			
└─└─┬─BatchNorm2d (1)	[32, 32, 112, 112]	[32, 32, 112, 112]	(64)
False			

		└─SiLU (2)	[32, 32, 112, 112]	[32, 32, 112, 112]	--
--					
		└─Sequential (1)	[32, 32, 112, 112]	[32, 16, 112, 112]	--
False					
		└─MBConv (0)	[32, 32, 112, 112]	[32, 16, 112, 112]	(1,448)
False					
		└─Sequential (2)	[32, 16, 112, 112]	[32, 24, 56, 56]	--
False					
		└─MBConv (0)	[32, 16, 112, 112]	[32, 24, 56, 56]	(6,004)
False					
		└─MBConv (1)	[32, 24, 56, 56]	[32, 24, 56, 56]	(10,710)
False					
		└─Sequential (3)	[32, 24, 56, 56]	[32, 40, 28, 28]	--
False					
		└─MBConv (0)	[32, 24, 56, 56]	[32, 40, 28, 28]	(15,350)
False					
		└─MBConv (1)	[32, 40, 28, 28]	[32, 40, 28, 28]	(31,290)
False					
		└─Sequential (4)	[32, 40, 28, 28]	[32, 80, 14, 14]	--
False					
		└─MBConv (0)	[32, 40, 28, 28]	[32, 80, 14, 14]	(37,130)
False					
		└─MBConv (1)	[32, 80, 14, 14]	[32, 80, 14, 14]	(102,900)
False					
		└─MBConv (2)	[32, 80, 14, 14]	[32, 80, 14, 14]	(102,900)
False					
		└─Sequential (5)	[32, 80, 14, 14]	[32, 112, 14, 14]	--
False					
		└─MBConv (0)	[32, 80, 14, 14]	[32, 112, 14, 14]	(126,004)
False					
		└─MBConv (1)	[32, 112, 14, 14]	[32, 112, 14, 14]	(208,572)
False					
		└─MBConv (2)	[32, 112, 14, 14]	[32, 112, 14, 14]	(208,572)
False					
		└─Sequential (6)	[32, 112, 14, 14]	[32, 192, 7, 7]	--
False					
		└─MBConv (0)	[32, 112, 14, 14]	[32, 192, 7, 7]	(262,492)
False					
		└─MBConv (1)	[32, 192, 7, 7]	[32, 192, 7, 7]	(587,952)
False					
		└─MBConv (2)	[32, 192, 7, 7]	[32, 192, 7, 7]	(587,952)
False					
		└─MBConv (3)	[32, 192, 7, 7]	[32, 192, 7, 7]	(587,952)
False					

	└─Sequential (7)	[32, 192, 7, 7]	[32, 320, 7, 7]	--
False				
	└─MBConv (0)	[32, 192, 7, 7]	[32, 320, 7, 7]	(717,232)
False				
	└─Conv2dNormActivation (8)	[32, 320, 7, 7]	[32, 1280, 7, 7]	--
False				
	└─Conv2d (0)	[32, 320, 7, 7]	[32, 1280, 7, 7]	(409,600)
False				
	└─BatchNorm2d (1)	[32, 1280, 7, 7]	[32, 1280, 7, 7]	(2,560)
False				
	└─SiLU (2)	[32, 1280, 7, 7]	[32, 1280, 7, 7]	--
--				
	└─AdaptiveAvgPool2d (avgpool)	[32, 1280, 7, 7]	[32, 1280, 1, 1]	--
--				
	└─Sequential (classifier)	[32, 1280]	[32, 10]	--
True				
	└─Dropout (0)	[32, 1280]	[32, 1280]	-- --
	└─Linear (1)	[32, 1280]	[32, 10]	12,810
True				

=====
 =====
 Total params: 4,020,358
 Trainable params: 12,810
 Non-trainable params: 4,007,548
 Total mult-adds (G): 12.31
 =====
 =====

Input size (MB): 19.27
 Forward/backward pass size (MB): 3452.09
 Params size (MB): 16.08
 Estimated Total Size (MB): 3487.44
 =====
 =====

Layer (type (var_name))	Input Shape	Output Shape	Param #
Trainable			
=====			
=====			
ResNet (ResNet)	[32, 3, 224, 224]	[32, 10]	-- Partial
└─Conv2d (conv1)	[32, 3, 224, 224]	[32, 64, 112, 112]	(9,408) False
└─BatchNorm2d (bn1)	[32, 64, 112, 112]	[32, 64, 112, 112]	(128) False

└─ReLU (relu)	[32, 64, 112, 112]	[32, 64, 112, 112]	--	--
└─MaxPool2d (maxpool)	[32, 64, 112, 112]	[32, 64, 56, 56]	--	--
└─Sequential (layer1)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─└─BasicBlock (0)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─└─└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─└─└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─└─BasicBlock (1)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─└─└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─└─└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─└─BasicBlock (2)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─└─└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─└─└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
└─└─└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
└─└─└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
└─Sequential (layer2)	[32, 64, 56, 56]	[32, 128, 28, 28]	--	False
└─└─BasicBlock (0)	[32, 64, 56, 56]	[32, 128, 28, 28]	--	False
└─└─└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 128, 28, 28]	(73,728)	False
└─└─└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	False
└─└─└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
└─└─└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	False
└─└─└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	False
└─└─Sequential (downsample)	[32, 64, 56, 56]	[32, 128, 28, 28]	(8,448)	False
└─└─└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
└─└─└─BasicBlock (1)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	False
└─└─└─└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	False
└─└─└─└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	False
└─└─└─└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--

False	└─Conv2d (conv2)	[32,128, 28, 28]	[32, 128, 28, 28]	(147,456)
False	└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False	└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--
False	└─BasicBlock (2)	[32, 128, 28, 28]	[32, 128, 28, 28]	False
False	└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)
False	└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False	└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--
False	└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)
False	└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False	└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--
False	└─BasicBlock (3)	[32, 128, 28, 28]	[32, 128, 28, 28]	False
False	└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)
False	└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False	└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--
False	└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)
False	└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False	└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--
False	└─Sequential (layer3)	[32, 128, 28, 28]	[32, 256, 14, 14]	False
False	└─BasicBlock (0)	[32, 128, 28, 28]	[32, 256, 14, 14]	False
False	└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 256, 14, 14]	(294,912)
False	└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False	└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--
False	└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False	└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False	└─Sequential (downsample)	[32, 128, 28, 28]	[32, 256, 14, 14]	(33,280)
False	└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--
False	└─BasicBlock (1)	[32, 256, 14, 14]	[32, 256, 14, 14]	False
False	└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)

		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (2)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (3)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (4)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (5)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					

		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		─Sequential (layer4)	[32, 256, 14, 14]	[32, 512, 7, 7]	--	False
		└─BasicBlock (0)	[32, 256, 14, 14]	[32, 512, 7, 7]	--	False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 512, 7, 7]	(1,179,648)	
False						
		└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─Sequential (downsample)	[32, 256, 14, 14]	[32, 512, 7, 7]	(132,096)	
False						
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─BasicBlock (1)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	False
		└─Conv2d (conv1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─BasicBlock (2)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	False
		└─Conv2d (conv1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		─AdaptiveAvgPool2d (avgpool)	[32, 512, 7, 7]	[32, 512, 1, 1]	--	--
		─Linear (fc)	[32, 512]	[32, 10]	5,130	True

Total params: 21,289,802
 Trainable params: 5,130
 Non-trainable params: 21,284,672
 Total mult-adds (G): 117.22

Input size (MB): 19.27

Forward/backward pass size (MB): 1913.92

Params size (MB): 85.16

Estimated Total Size (MB): 2018.34

=====

- Training (fine tuning the models for the same) SVHN.
- Training Results for both:

Effnet:

Epoch: 1 | train_loss: 1.7303 | train_acc: 0.3565 | test_loss: 1.4631 | test_acc: 0.4466
Epoch: 2 | train_loss: 1.6190 | train_acc: 0.4052 | test_loss: 1.4546 | test_acc: 0.4525
Epoch: 3 | train_loss: 1.6110 | train_acc: 0.4099 | test_loss: 1.4096 | test_acc: 0.4650
Epoch: 4 | train_loss: 1.6074 | train_acc: 0.4114 | test_loss: 1.4169 | test_acc: 0.4521
Epoch: 5 | train_loss: 1.6020 | train_acc: 0.4132 | test_loss: 1.4224 | test_acc: 0.4500
Epoch: 6 | train_loss: 1.6007 | train_acc: 0.4158 | test_loss: 1.4077 | test_acc: 0.4611
Epoch: 7 | train_loss: 1.6063 | train_acc: 0.4113 | test_loss: 1.3802 | test_acc: 0.4613
Epoch: 8 | train_loss: 1.6057 | train_acc: 0.4110 | test_loss: 1.4265 | test_acc: 0.4570
Epoch: 9 | train_loss: 1.6013 | train_acc: 0.4137 | test_loss: 1.4317 | test_acc: 0.4498
Epoch: 10 | train_loss: 1.6027 | train_acc: 0.4115 | test_loss: 1.3979 | test_acc: 0.4638
Epoch: 11 | train_loss: 1.6032 | train_acc: 0.4131 | test_loss: 1.4110 | test_acc: 0.4630
Epoch: 12 | train_loss: 1.6024 | train_acc: 0.4127 | test_loss: 1.3908 | test_acc: 0.4544

ResNet:



Epoch: 1 | train_loss: 1.8573 | train_acc: 0.2952 | test_loss: 1.7167 | test_acc: 0.3452
Epoch: 2 | train_loss: 1.7235 | train_acc: 0.3520 | test_loss: 1.6756 | test_acc: 0.3541
Epoch: 3 | train_loss: 1.6972 | train_acc: 0.3673 | test_loss: 1.6537 | test_acc: 0.3658
Epoch: 4 | train_loss: 1.6818 | train_acc: 0.3702 | test_loss: 1.6454 | test_acc: 0.3716
Epoch: 5 | train_loss: 1.6712 | train_acc: 0.3762 | test_loss: 1.6442 | test_acc: 0.3835
Epoch: 6 | train_loss: 1.6694 | train_acc: 0.3769 | test_loss: 1.6279 | test_acc: 0.3645
Epoch: 7 | train_loss: 1.6589 | train_acc: 0.3835 | test_loss: 1.6550 | test_acc: 0.3609
Epoch: 8 | train_loss: 1.6581 | train_acc: 0.3827 | test_loss: 1.6174 | test_acc: 0.3942
Epoch: 9 | train_loss: 1.6576 | train_acc: 0.3843 | test_loss: 1.6180 | test_acc: 0.3659
Epoch: 10 | train_loss: 1.6505 | train_acc: 0.3860 | test_loss: 1.6000 | test_acc: 0.3761
Epoch: 11 | train_loss: 1.6520 | train_acc: 0.3853 | test_loss: 1.6444 | test_acc: 0.3671
Epoch: 12 | train_loss: 1.6504 | train_acc: 0.3873 | test_loss: 1.6192 | test_acc: 0.3867
Epoch: 13 | train_loss: 1.6536 | train_acc: 0.3863 | test_loss: 1.5860 | test_acc: 0.3891
Epoch: 14 | train_loss: 1.6505 | train_acc: 0.3853 | test_loss: 1.6426 | test_acc: 0.3721
Epoch: 15 | train_loss: 1.6458 | train_acc: 0.3879 | test_loss: 1.6095 | test_acc: 0.3886
total training time: 1414.568 sec.

- Saving the trained model, loading, and re-evaluating the training.
- Testing the Impact of Attacks on the models
- Defining attacks effnet PGD:
Purtubed Test accuracy: 0.0864
- Defining attacks resnet PGD:
Purtubed Test accuracy: 0.0183
- Defining attacks effnet Jitter:
Purtubed Test accuracy: 0.0721
- Defining attacks resnet Jitter:
Purtubed Test accuracy: 0.0362

Procedure For task 3 (detect adversarial attacks):

- Import required packages. (PyTorch, NumPy, Matplotlib, torchmetrics ...etc.)
- Setting up device-agnostic code for faster training.
- Downloading the Dataset using PyTorch vision datasets.
- Modify the dataset to make 2 classes, attacked or non-attacked data.
- Modification is done with help of the previous saved model.
- Train the model for 2 class classification for the above dataset.
- Train Results:

Epoch: 1 | train_loss: 0.5410 | train_acc: 0.7240 | test_loss: 0.5188 | test_acc: 0.7564
 Epoch: 2 | train_loss: 0.4856 | train_acc: 0.7679 | test_loss: 0.5127 | test_acc: 0.7582
 Epoch: 3 | train_loss: 0.4766 | train_acc: 0.7736 | test_loss: 0.4988 | test_acc: 0.7667
 Epoch: 4 | train_loss: 0.4792 | train_acc: 0.7775 | test_loss: 0.4698 | test_acc: 0.7836
 Epoch: 5 | train_loss: 0.4928 | train_acc: 0.7599 | test_loss: 0.4978 | test_acc: 0.7711
 total training time: 75.495 sec.



- Checking FGSM with different epsilons.

Perturbed Test accuracy: 0.7739 with epsilon: 0.0
Perturbed Test accuracy: 0.7751 with epsilon: 0.0625
Perturbed Test accuracy: 0.7705 with epsilon: 0.125
Perturbed Test accuracy: 0.7643 with epsilon: 0.1875
Perturbed Test accuracy: 0.7665 with epsilon: 0.25
Perturbed Test accuracy: 0.7759 with epsilon: 0.3125
Perturbed Test accuracy: 0.7691 with epsilon: 0.375
Perturbed Test accuracy: 0.7598 with epsilon: 0.4375
Perturbed Test accuracy: 0.7785 with epsilon: 0.5
Perturbed Test accuracy: 0.7657 with epsilon: 0.5625
Perturbed Test accuracy: 0.7665 with epsilon: 0.625
Perturbed Test accuracy: 0.7669 with epsilon: 0.6875
Perturbed Test accuracy: 0.7729 with epsilon: 0.75
Perturbed Test accuracy: 0.7757 with epsilon: 0.8125
Perturbed Test accuracy: 0.7701 with epsilon: 0.875
Perturbed Test accuracy: 0.7703 with epsilon: 0.9375

- Doing same with saving and loading model for test.

Perturbed Test accuracy: 0.7739 with epsilon: 0.0
Perturbed Test accuracy: 0.7751 with epsilon: 0.0625
Perturbed Test accuracy: 0.7705 with epsilon: 0.125
Perturbed Test accuracy: 0.7643 with epsilon: 0.1875
Perturbed Test accuracy: 0.7665 with epsilon: 0.25
Perturbed Test accuracy: 0.7759 with epsilon: 0.3125
Perturbed Test accuracy: 0.7691 with epsilon: 0.375
Perturbed Test accuracy: 0.7598 with epsilon: 0.4375
Perturbed Test accuracy: 0.7785 with epsilon: 0.5
Perturbed Test accuracy: 0.7657 with epsilon: 0.5625
Perturbed Test accuracy: 0.7665 with epsilon: 0.625
Perturbed Test accuracy: 0.7669 with epsilon: 0.6875
Perturbed Test accuracy: 0.7729 with epsilon: 0.75
Perturbed Test accuracy: 0.7757 with epsilon: 0.8125
Perturbed Test accuracy: 0.7701 with epsilon: 0.875
Perturbed Test accuracy: 0.7703 with epsilon: 0.9375

Task: Deepfake Detection

Objectives:

- Create 100 deepfakes/faceswap from the existing tools of your choice using your own face images.
- Split this data into fine-tuned and test sets (50-50).
- Finetune your model in Q1(iii) and test on the remaining 50 test samples, and report the performance.

Procedure:

- Import required packages. (PyTorch, NumPy, Matplotlib, torchmetrics ...etc.)
- Setting up device-agnostic code for faster training.
- Loading Saved model from question 1, part 3.
- Model Summary:

```
=====
```

Layer (type (var_name))	Input Shape	Output Shape	Param #
Trainable			
=====			
=====			
ResNet (ResNet)	[32, 3, 224, 224]	[32, 2]	-- Partial
└─Conv2d (conv1)	[32, 3, 224, 224]	[32, 64, 112, 112]	(9,408) False
└─BatchNorm2d (bn1)	[32, 64, 112, 112]	[32, 64, 112, 112]	(128) False
└─ReLU (relu)	[32, 64, 112, 112]	[32, 64, 112, 112]	-- --
└─MaxPool2d (maxpool)	[32, 64, 112, 112]	[32, 64, 56, 56]	-- --
└─Sequential (layer1)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- False
└─BasicBlock (0)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- False
└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False
└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128) False
└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- --
└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False
└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128) False
└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- --
└─BasicBlock (1)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- False
└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False
└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128) False
└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- --
└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False
└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128) False
└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- --
└─BasicBlock (2)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- False
└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False
└─BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128) False
└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	-- --
└─Conv2d (conv2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864) False

		└─BatchNorm2d (bn2)	[32, 64, 56, 56]	[32, 64, 56, 56]	(128)	False
		└─ReLU (relu)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	--
	└─Sequential (layer2)		[32, 64, 56, 56]	[32, 128, 28, 28]	--	False
		└─BasicBlock (0)	[32, 64, 56, 56]	[32, 128, 28, 28]	--	False
		└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 128, 28, 28]	(73,728)	False
		└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─Sequential (downsample)	[32, 64, 56, 56]	[32, 128, 28, 28]	(8,448)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─BasicBlock (1)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	False
		└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─BasicBlock (2)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	False
		└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--
		└─BasicBlock (3)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	False
		└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)	
False						
		└─BatchNorm2d (bn1)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)	
False						
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	--	--

		└─Conv2d (conv2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(147,456)
False					
		└─BatchNorm2d (bn2)	[32, 128, 28, 28]	[32, 128, 28, 28]	(256)
False					
		└─ReLU (relu)	[32, 128, 28, 28]	[32, 128, 28, 28]	-- --
		└─Sequential (layer3)	[32, 128, 28, 28]	[32, 256, 14, 14]	-- False
		└─BasicBlock (0)	[32, 128, 28, 28]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 128, 28, 28]	[32, 256, 14, 14]	(294,912)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─Sequential (downsample)	[32, 128, 28, 28]	[32, 256, 14, 14]	(33,280)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (1)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (2)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)
False					
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- --
		└─BasicBlock (3)	[32, 256, 14, 14]	[32, 256, 14, 14]	-- False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)
False					

		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─BasicBlock (4)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─BasicBlock (5)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn1)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─Conv2d (conv2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(589,824)	
False						
		└─BatchNorm2d (bn2)	[32, 256, 14, 14]	[32, 256, 14, 14]	(512)	
False						
		└─ReLU (relu)	[32, 256, 14, 14]	[32, 256, 14, 14]	--	--
		└─Sequential (layer4)	[32, 256, 14, 14]	[32, 512, 7, 7]	--	False
		└─BasicBlock (0)	[32, 256, 14, 14]	[32, 512, 7, 7]	--	False
		└─Conv2d (conv1)	[32, 256, 14, 14]	[32, 512, 7, 7]	(1,179,648)	
False						
		└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False
		└─Sequential (downsample)	[32, 256, 14, 14]	[32, 512, 7, 7]	(132,096)	
False						
		└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--
		└─BasicBlock (1)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	False
		└─Conv2d (conv1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)	False
		└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)	False

└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--	
└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)		False
└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)		False
└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--	
└─BasicBlock (2)	[32, 512, 7, 7]	[32, 512, 7, 7]	--		False
└─Conv2d (conv1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)		False
└─BatchNorm2d (bn1)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)		False
└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--	
└─Conv2d (conv2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(2,359,296)		False
└─BatchNorm2d (bn2)	[32, 512, 7, 7]	[32, 512, 7, 7]	(1,024)		False
└─ReLU (relu)	[32, 512, 7, 7]	[32, 512, 7, 7]	--	--	
└─AdaptiveAvgPool2d (avgpool)	[32, 512, 7, 7]	[32, 512, 1, 1]	--		--
└─Linear (fc)	[32, 512]	[32, 2]	1,026		True

Total params: 21,285,698

Trainable params: 1,026

Non-trainable params: 21,284,672

Total mult-adds (G): 117.22

Input size (MB): 19.27

Forward/backward pass size (MB): 1913.91

Params size (MB): 85.14

Estimated Total Size (MB): 2018.32

- Loading Created Dataset with the help of torchvision.datasets.ImageFolder
- Dataset Info:

(Dataset ImageFolder

Number of datapoints: 108

Root location: data/gan_custom_data/train/

StandardTransform

Transform: ImageClassification(

crop_size=[224]

resize_size=[256]

mean=[0.485, 0.456, 0.406]

std=[0.229, 0.224, 0.225]

interpolation=InterpolationMode.BILINEAR

),

Dataset ImageFolder

Number of datapoints: 108

Root location: data/gan_custom_data/test/

StandardTransform


```
Transform: ImageClassification(  
    crop_size=[224]  
    resize_size=[256]  
    mean=[0.485, 0.456, 0.406]  
    std=[0.229, 0.224, 0.225]  
    interpolation=InterpolationMode.BILINEAR  
)
```

- Visualizing random datasamples



- Here Label 0 means an original image with applied augmentation, while 1 Represents modified/ Gan Generated.
- Converting the dataset to data loader.
- Training the model for the same
- Training Results:

Epoch: 1 | train_loss: 0.7352 | train_acc: 0.5026 | test_loss: 0.5626 | test_acc: 0.7188
 Epoch: 2 | train_loss: 0.5355 | train_acc: 0.8177 | test_loss: 0.4434 | test_acc: 0.8906
 Epoch: 3 | train_loss: 0.3730 | train_acc: 0.9245 | test_loss: 0.3877 | test_acc: 0.9219
 Epoch: 4 | train_loss: 0.2758 | train_acc: 0.9609 | test_loss: 0.3398 | test_acc: 0.9453
 Epoch: 5 | train_loss: 0.2213 | train_acc: 0.9688 | test_loss: 0.3102 | test_acc: 0.9453
 Epoch: 6 | train_loss: 0.1904 | train_acc: 0.9688 | test_loss: 0.2904 | test_acc: 0.9688
 Epoch: 7 | train_loss: 0.1680 | train_acc: 0.9844 | test_loss: 0.3070 | test_acc: 0.8724
 Epoch: 8 | train_loss: 0.1502 | train_acc: 0.9635 | test_loss: 0.2753 | test_acc: 0.9349
 Epoch: 9 | train_loss: 0.1219 | train_acc: 0.9922 | test_loss: 0.2611 | test_acc: 0.9635
 Epoch: 10 | train_loss: 0.0978 | train_acc: 0.9922 | test_loss: 0.2794 | test_acc: 0.8932
 Epoch: 11 | train_loss: 0.0878 | train_acc: 0.9922 | test_loss: 0.2608 | test_acc: 0.9505
 Epoch: 12 | train_loss: 0.0846 | train_acc: 0.9792 | test_loss: 0.2684 | test_acc: 0.9219
 Epoch: 13 | train_loss: 0.0935 | train_acc: 0.9922 | test_loss: 0.2978 | test_acc: 0.8333
 Epoch: 14 | train_loss: 0.0690 | train_acc: 1.0000 | test_loss: 0.2654 | test_acc: 0.8724
 Epoch: 15 | train_loss: 0.0645 | train_acc: 0.9922 | test_loss: 0.2473 | test_acc: 0.9505
 total training time: 43.407 sec.



- Doing the same with saving and loading models for testing purposes.
 test_loss: 0.2473 | test_acc: 0.9505

Task: Audio Deepfake Detection

Objectives:

- Record 1000 hindi and 1000 english sentences in your voice sampled from the given text files
- generated audios for spoof detection i.e real vs fake classification.
- Share the recorded and generated audios within separate folders respectively, using the drive link.

Procedure:

- Record using voice:
 - https://www.dictate.app/annotation/record/selected_lines_eng_20230221_171133#
 - https://www.dictate.app/annotation/record/selected_lines_hin_20230224_072858#
- Generated Voice Using Google Text To Speech Engine:

Python

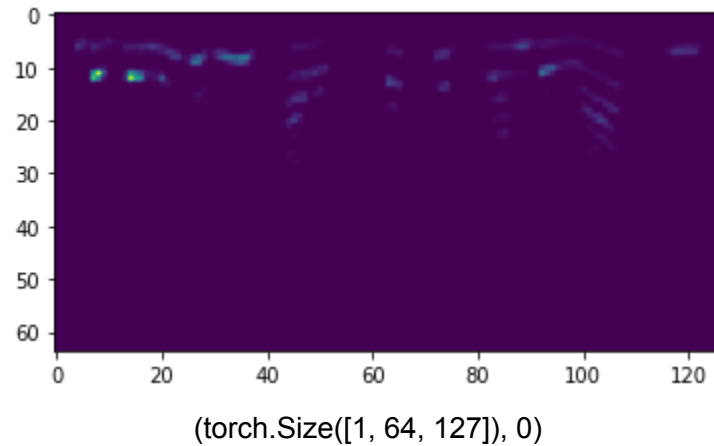
```
from gtts import gTTS
from pathlib import Path

language = 'en'
# language = 'hi'

# mytext = 'this is an example in english!'
# mytext = 'यह हिंदी में एक उदाहरण है'
Path("data/english").mkdir(parents=True, exist_ok=True)
with open("selected_lines_eng.txt") as fp:
    lines = fp.readlines()
    i=0
    for mytext in lines[:5]:
        myobj = gTTS(text=mytext, lang=language,
tld='co.in', slow=False)
        myobj.save(f"data/english/{i}_female_eng.mp3")
        print(f"{i}/{len(lines)} done...")
        i += 1
```

- Restructuring the dataset.
- Dataset Link:-
https://drive.google.com/file/d/1T0BpspJHZ9aEDztdvVc9T5bsblZITnz_/view?usp=share_link
- Generating CSV File for the dataset will be used in Custom DATASET Class.
- Defining Classes for the Dataset - Spectrogram

- One sample from the dataset.



- Creating Data loaders with 80-20 Split of train and test data loaders.
- Defining Training and Testing Steps function for Training the Model.
- Defining models with pre-trained weights.
- Model Architecture:

=====				
=====				
Layer (type (var_name))	Input Shape	Output Shape	Param #	
Trainable				
=====				
=====				
VGG (VGG)	[32, 1, 64, 173]	[32, 2]	--	Partial
└─Sequential (features)	[32, 1, 64, 173]	[32, 512, 2, 5]	--	False
└─┬─Conv2d (0)	[32, 1, 64, 173]	[32, 64, 64, 173]	(640)	False
└─┬─ReLU (1)	[32, 64, 64, 173]	[32, 64, 64, 173]	--	--
└─┬─MaxPool2d (2)	[32, 64, 64, 173]	[32, 64, 32, 86]	--	--
└─┬─Conv2d (3)	[32, 64, 32, 86]	[32, 128, 32, 86]	(73,856)	False
└─┬─ReLU (4)	[32, 128, 32, 86]	[32, 128, 32, 86]	--	--
└─┬─MaxPool2d (5)	[32, 128, 32, 86]	[32, 128, 16, 43]	--	--
└─┬─Conv2d (6)	[32, 128, 16, 43]	[32, 256, 16, 43]	(295,168)	False
└─┬─ReLU (7)	[32, 256, 16, 43]	[32, 256, 16, 43]	--	--
└─┬─Conv2d (8)	[32, 256, 16, 43]	[32, 256, 16, 43]	(590,080)	False
└─┬─ReLU (9)	[32, 256, 16, 43]	[32, 256, 16, 43]	--	--
└─┬─MaxPool2d (10)	[32, 256, 16, 43]	[32, 256, 8, 21]	--	--
└─┬─Conv2d (11)	[32, 256, 8, 21]	[32, 512, 8, 21]	(1,180,160)	False
└─┬─ReLU (12)	[32, 512, 8, 21]	[32, 512, 8, 21]	--	--
└─┬─Conv2d (13)	[32, 512, 8, 21]	[32, 512, 8, 21]	(2,359,808)	False
└─┬─ReLU (14)	[32, 512, 8, 21]	[32, 512, 8, 21]	--	--
└─┬─MaxPool2d (15)	[32, 512, 8, 21]	[32, 512, 4, 10]	--	--
└─┬─Conv2d (16)	[32, 512, 4, 10]	[32, 512, 4, 10]	(2,359,808)	False
└─┬─ReLU (17)	[32, 512, 4, 10]	[32, 512, 4, 10]	--	--

└─Conv2d (18)	[32, 512, 4, 10]	[32, 512, 4, 10]	(2,359,808)	False
└─ReLU (19)	[32, 512, 4, 10]	[32, 512, 4, 10]	--	--
└─MaxPool2d (20)	[32, 512, 4, 10]	[32, 512, 2, 5]	--	--
└─AdaptiveAvgPool2d (avgpool)	[32, 512, 2, 5]	[32, 512, 7, 7]	--	--
└─Sequential (classifier)	[32, 25088]	[32, 2]	--	True
└─Linear (0)	[32, 25088]	[32, 4096]	102,764,544	True
└─ReLU (1)	[32, 4096]	[32, 4096]	--	--
└─Dropout (2)	[32, 4096]	[32, 4096]	--	--
└─Linear (3)	[32, 4096]	[32, 4096]	16,781,312	True
└─ReLU (4)	[32, 4096]	[32, 4096]	--	--
└─Dropout (5)	[32, 4096]	[32, 4096]	--	--
└─Linear (6)	[32, 4096]	[32, 2]	8,194	True

=====
Total params: 128,773,378

Trainable params: 119,554,050

Non-trainable params: 9,219,328

Total mult-adds (G): 55.12
=====

=====
Input size (MB): 1.42

Forward/backward pass size (MB): 418.38

Params size (MB): 515.09

Estimated Total Size (MB): 934.89
=====

• Trained Model Results;

Epoch: 1 | train_loss: 53.9074 | train_acc: 0.9081 | test_loss: 42.5900 | test_acc: 0.9075

Epoch: 2 | train_loss: 20.1229 | train_acc: 0.9522 | test_loss: 4.0608 | test_acc: 0.9738

Epoch: 3 | train_loss: 3.9477 | train_acc: 0.9759 | test_loss: 0.9426 | test_acc: 0.9938

Epoch: 4 | train_loss: 25.0096 | train_acc: 0.9628 | test_loss: 50.8121 | test_acc: 0.9750

Epoch: 5 | train_loss: 86.8253 | train_acc: 0.9544 | test_loss: 16.2660 | test_acc: 0.9938

Epoch: 6 | train_loss: 27.7979 | train_acc: 0.9822 | test_loss: 1.1427 | test_acc: 0.9962

Epoch: 7 | train_loss: 10.8485 | train_acc: 0.9850 | test_loss: 3.0007 | test_acc: 0.9888

Epoch: 8 | train_loss: 5.0821 | train_acc: 0.9888 | test_loss: 19.8806 | test_acc: 0.9688

Epoch: 9 | train_loss: 3.6336 | train_acc: 0.9866 | test_loss: 1.4706 | test_acc: 0.9975

Epoch: 10 | train_loss: 1.4425 | train_acc: 0.9931 | test_loss: 1.7719 | test_acc: 0.9888

Trained feed forward net saved at model_pre_trained.pth

• Training Other CNN Arch. for the same

- Results:

Epoch: 1 | train_loss: 0.8077 | train_acc: 0.5056 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 2 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 3 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 4 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 5 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 6 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 7 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 8 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 9 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Epoch: 10 | train_loss: 0.8079 | train_acc: 0.5053 | test_loss: 0.8095 | test_acc: 0.5038
Trained feed forward net saved at model_custom_cnn_arch.pth

ALL CODE + DATA:- [A1](#)

Note:-

- Pre Trained Arch. performs very well because it is pre-trained, and the dataset is very small for training a network from scratch.
- Please use the IITJ Email id for accessing the data and the code.

References:

<https://arxiv.org/abs/2010.01950>
<https://arxiv.org/pdf/2110.01200.pdf>
<https://github.com/clovaai/aasist>
<https://github.com/Harry24k/adversarial-attacks-pytorch>
<https://adversarial-attacks-pytorch.readthedocs.io/en/latest/attacks.html>
<https://replicate.com/yoyo-nb/thin-plate-spline-motion-model>
<https://replicate.com/yoyo-nb/thin-plate-spline-motion-model/api>
<https://giphy.com/> -> For deep fake samples references
<https://gtts.readthedocs.io/en/latest/>
<https://www.dictate.app/> -> For Recoding Audio
<https://pytorch.org/docs/stable/index.html>
<https://numpy.org/doc/stable/reference/index.html>
<https://torchmetrics.readthedocs.io/en/stable/>
<https://stackoverflow.com/>